

A benchmark suite and performance analysis of provenance systems

Samuel Grayson

grayson5@illinois.edu

University of Illinois Urbana Champaign

Department of Computer Science

Urbana, IL, USA

Daniel S. Katz

dskatz@illinois.edu

University of Illinois Urbana Champaign

Department of Computer Science

National Center for Supercomputing Applications

Department of Electrical and Computer Engineering

School of Information Sciences

Urbana, IL, USA

Reed Milewicz

rmilewi@sandia.gov

Sandia National Laboratories

Software Engineering and Research Department

Albuquerque, NM, USA

Darko Marinov

marinov@illinois.edu

University of Illinois Urbana Champaign

Department of Computer Science

Urbana, IL, USA

ACM Reference Format:

Samuel Grayson, Daniel S. Katz, Reed Milewicz, and Darko Marinov. 2024. A benchmark suite and performance analysis of provenance systems. In . ACM, New York, NY, USA, 9 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 BACKGROUND

Computational provenance, “the computational input artifacts and computational processes that influenced a certain computational output artifact” [20], has many potential applications, including the following from Pimentel et al. [44] and Sar and Cao [47]:

1. **Reproducibility.** A description of the inputs and processes used to generate a specific output can aid manual and automatic reproduction of that output¹. Lack of reproducibility in computational experiments undermines the long-term credibility of science and hinders the day-to-day work of researchers. Empirical studies [17, 23, 54, 58] show that reproducibility is rarely achieved in practice, probably due to its difficulty under the short time budget that scientists have available to spend on reproducibility. If reproducibility was easier to attain, perhaps because of automatic provenance tracking, it may improve the reproducibility rate of computational research.
- **Manual reproducibility.** A user can Provenance data improves manual reproducibility, because users have a record of the inputs, outputs, and processes used to create a computational artifact.

¹“Reproduction”, in the ACM sense, where a **different team** uses the **same artifacts** to generate the output artifact [6].

Unpublished working draft. Not for distribution.

Permission to make digital or hard copies of all or part of this work for personal or internal use, or the internal or personal use of specific clients, is granted by ACM for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Conference’17, July 2017, Washington, DC, USA

© 2024 Association for Computing Machinery.

ACM ISBN 978-x-xxxx-xxxx-x/YY/MM...\$15.00

<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

2024-02-08 04:34. Page 1 of 1–9.

- **Automatic reproducibility.** Provenance data also has the potential to enable automatic reproducibility, if the process trace is detailed enough to be “re-executed”. This idea is also called “software record/replay”. However, not all provenance systems make this their goal.
2. **Caching subsequent re-executions.** Computational science inquiries involves changing some code and re-executing the workflows (e.g., testing different clustering algorithms). In these cases, the user has to keep track of what parts of the code they changed, and which process have to be re-executed. However, an automated system could read the computational provenance graphs produced by previous executions, look at what parts of the code changed, and safely decide what processes need to be re-executed. The dependency graph would be automatically deduced, leaving less chance for a dependency-misspecification, unlike Make or CMake, which requires the user to manually specify a dependency graph.
3. **Comprehension.** Provenance helps the user understand and document workflows. An automated tool that consumes provenance can answer queries like “What version of the data did I use for this figure?” and “Does this workflow include FERPA-protected data?”.
4. **Data cataloging.** Provenance data can help catalog, label, and recall experimental results based on the input parameters. For example, a user might have run dozens of different versions of their workflow, and they may want to ask an automated system, “show me the results I previously computed based on that data with this algorithm?”.
5. **Space compression.** If the provenance of a particular artifact is known, the artifact may be able to be deleted to save space, and regenerated by need.

One can capture computational provenance by modifying an application to report provenance data, leveraging a workflow engine or programming language to report provenance data, or leveraging an operating system to emit provenance data to report provenance data [20].

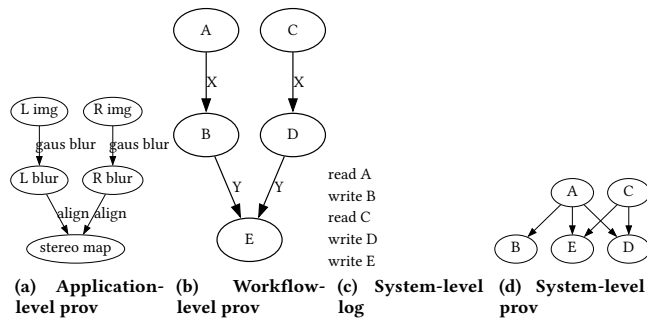


Figure 1: Several provenance graphs collected at different levels.

- **Application-level** provenance is the most semantically rich, since it knows the use of each input at the application-level (see fig. 1a), but the least general, since each application would have to be modified individually.
- **Workflow-level or language-level** provenance is a middle ground in semantic richness and generality; it only knows the use of inputs in a dataflow sense (see fig. 1b), but all applications using the provenance-modified workflow engine or programming language would emit provenance data without themselves being modified to emit provenance data.
- **System-level** is the most general, since all applications on the system would emit provenance data, but it is the least semantically rich, since observed dependencies may overapproximate the true dependencies (see fig. 1c and fig. 1d).

One may imagine an abstract tradeoff curve (??) between “enabling provenance applications such as reproducibility” as the horizontal axis increasing rightwards and “cost of implementation” that provenance data on the vertical axis increasing upwards). A typical status quo, not collecting any provenance data and not using workflows, is at the bottom left: no added cost and does nothing to enable provenance applications. System-level, workflow/language-level, and application-level are on a curve, increasing cost and enabling more provenance applications.

The initial jump in cost from nothing to system-level is low because the user need not change *anything* about their application; they merely need to install some provenance tracer onto their system and run their code, without modifying it, in the tracer.² Perceived ease of use is a critical factor in the adoption of new technologies (formalized in the Technology Acceptance Model [18]). Although the user may eventually use more semantically rich provenance, low-cost system-level provenance would get provenance’s “foot in the door”. While this data is less rich than that of the workflow or application level, it may be enough to enable important applications such as reproducibility, caching, etc. Since system-level provenance collection is a possibly valuable tradeoff between cost

²DSK: what about the performance penalty? Since you talk about performance in contributions, I think you have to introduce it here. SAG: This is referring to the “cost of switching from no-prov to prov”, which is low, and I’m only using this argument to explain why I look at system-level over the others. Performance overhead between system-level tools is a concern that I will address later on.

and enabling provenance applications, system-level provenance will be the subject of this work.

While there is little cost human-overhead in using system-level provenance (no user code change), there is a non-trivial implicit overhead in monitoring and recording each computational process. Even a minor overhead per I/O operation would get amplified over the tens of thousands of I/O operations that a program might execute per second.

Prior work in system-level provenance usually contains some benchmark programs to evaluate the overhead imposed by the system-level provenance tool. However, the set of chosen benchmark programs are not consistent from one publication to another, and overhead can be extremely sensitive to the exact choice of program, so results are totally incomparable between publications. Most publications only benchmark their new system against native/no-provenance, so prior work cannot easily establish which system-level provenance tool is the fastest.

2 CONTRIBUTIONS

This work aims to summarize state of the art, establish goalposts for future research in the area, and identify which provenance tools are practically usable.

This work contributes:

- **A rapid review:** There are scores of academic publications on system-level provenance (see table 2), so we collate as many provenance tools as possible and classify them by *capture method* (e.g., does the provenance collector require you to load a kernel module or run your code in a VM?).
- **A benchmark suite:** Prior work does not use a consistent set of benchmarks; often publications use an overlapping set of benchmarks from prior work. We collate benchmarks used in prior work, add some unrepresented areas, and find a statistically valid subset of the benchmark.
- **A quantitative performance comparison:** Prior publications often only compares the performance their provenance tool to the baseline, no-provenance performance, not to other provenance tools. It is difficult to compare provenance tools, given data of different benchmarks on different machines. This work runs a consistent set of benchmarks on a single machine over all provenance tools.
- **A predictive performance model:** The performance overhead of a single provenance system varies from <1% to 23% [40] based on the application, so a single number for overhead is not sufficient. This work develops a statistical model for predicting the overhead of \$X application in \$Y provenance system based on \$Y provenance system’s performance on our benchmark suite and \$X application’s performance characteristics (e.g., number of I/O syscalls).

3 METHODS

3.1 Rapid review

We began a rapid review to identify the research state-of-the-art tools for automatic system-level provenance.

Rapid Reviews are a lighter-weight alternative to systematic literature reviews with a focus on timely feedback for decision-making. Schünemann and Moja [48] show that Rapid Reviews can

yield substantially similar results to a systematic literature review, albeit with less detail. Although developed in medicine, Cartaxo et al. show that Rapid Reviews are useful for informing software engineering design decisions [13, 14].

We conducted a rapid review with the following parameters:

- **Objective:** Identify system-level provenance collection tools.
- **Search terms:** “system-level” AND “provenance”
- **Search engine:** Google Scholar
- **Number of results:** 50
 - This threshold is the point of diminishing returns, as no new tools came up in the 40th – 50th results.
- **Criteria:** A relevant publication would center on one or more operating system-level tools that capture file provenance. A tool requiring that the user use a specific application or platform would be irrelevant.

We record the following features for each system-level provenance tool:

- **Capture method:** What method does the tool use to capture provenance?
 - **User-level tracing:** A provenance tool may use “debugging” or “tracing” features provided by the kernel to trace the I/O operations of another program, e.g., `ptrace(2)` [4].
 - **Built-in auditing service:** A provenance tool may use auditing service built in to the kernel, e.g., Linux Auditing Framework [38], enhanced Berkeley Packet Filter (eBPF) [2], kprobes [30], and ETW [5] for Windows.
 - **Filesystem instrumentation:** A provenance tool may set up a file system, so it can log I/O operations, e.g., using Filesystem in User Space (FUSE) interface [3], or Virtual File System (VFS) interface [22].
 - **Dynamic library instrumentation:** A provenance tool may replace a library used to execute I/O operations (e.g., glibc) with one that logs the calls before executing them.
 - **Binary instrumentation:** A provenance tool may use binary instrumentation (dynamic or static) to identify I/O operations in another program.
 - **Compile-time instrumentation:** A provenance tool may be a compiler pass that modifies the program to emit provenance data, especially intra-program control flow.
 - **Kernel instrumentation:** A provenance tool may be a modified kernel either by directly modifying and recompiling the kernel’s source tree.
 - **Kernel module:** Rather than directly modify the kernel’s source, the provenance tool may simply require that the user load a custom kernel module.
 - **VM instrumentation:** A provenance tool may execute the program in a virtual machine, where it can observe the program’s I/O operations.

3.2 Benchmark selection

Using the tools selected above, we identified all benchmarks that have been used in prior work. We excluded benchmarks for which we could not even find the original program (e.g., TextTransfer), benchmarks that were not available for Linux (e.g., Internet Explorer), benchmarks with a graphical component (e.g., Notepad++),

Table 1: Our experimental machine description.

Name	Value
CPU	11th Gen Intel(R) Core(TM) i7-1165G7 @ 2.80GHz
RAM	16 GiB of SODIMM DDR4 Synchronous 2400 MHz
Kernel	Linux 6.1.64

or benchmarks with an interactive component (e.g., GNU Midnight Commander).

We implemented the benchmarks as packages for the Nix package manager³, so they are runnable on many different platforms. Nix has official installers for Linux, Mac OS X, and Windows Subsystem for Linux on i686, x86_64, and aarch64 architectures, but FreeBSD and OpenBSD both package Nix themselves, and it can likely be built from source on even more platforms.

We also added new benchmarks:

- **Data science:** None of the benchmarks resembled a typical data science program, so we added the most popular Notebooks from Kaggle.com, a data science competition website.
- **Compilations:** Prior work uses compilation of Apache or of Linux. We added compilation of several other packages (any package in Spack) to our benchmark. Compiling packages is a good use-case for a provenance tracer, because a user might trial-and-error multiple compile commands and not remember the exact sequence of “correct” commands; the provenance tracker would be able to recall the commands which did not get overwritten, so the user can know what commands “actually worked”.⁴

3.3 Performance experiment

To get consistent measurements, we select as many benchmarks and provenance tracers as we reasonably can, and run a complete matrix (every tracer on every benchmark). Table 1 describes our experimental machine. We use BenchExec [11] to precisely measure the CPU time, wall time, memory utilization, and other attributes of the process (including child processes) in a Linux CGroup without networking, isolated from other processes on the system.

3.4 Benchmark subsetting

We implemented and ran many different benchmarks, which may be costly for future researchers seeking to evaluate new provenance systems. Given the less-costly results of a small number of benchmarks, perhaps one predict the performance of the rest of the benchmarks. We will try several methods to identify the most important benchmarks, keeping the ones which perform the best:

- **Principal component analysis (PCA) and K-means.** This is the traditional benchmark subsetting procedure evaluated by Yi et al. [?].
 1. We form a matrix of all benchmarks by “observed features” of that benchmark. The observed features contain, for example, how many file open/closes per second, how many file metadata operations per second, etc. While not traditionally done, “the overhead ratio of this benchmark in that provenance engine” is a perfectly valid feature to consider.

³See <https://nixos.org/guides/how-nix-works>

⁴DSK: this reminds me of VisTrails from Utah

We only use features that are invariant when running the program twice (i.e., using operations per second, not total operations), so benchmarks which are functionally similar but run for a longer time to not consist different points than analogous shorter benchmarks.

2. We apply PCA to that matrix. PCA is a mathematical procedure which combines a large number of “observed features” into smaller number “virtual features”, linearly, while maximizing the amount of variance in the resulting “virtual space” (in a sense, spreading out the benchmarks as much as possible from each other).
3. We apply K-means to the benchmarks in their reduced PCA-space. K-means is a fast clustering algorithm. Once the benchmarks are grouped into clusters, we identify one benchmark from each cluster to consist the benchmark subset.

- **Interpolative decomposition (ID).**

1. We form a matrix where each benchmark is a column, each provenance system is a row, and the elements contain the log of the overhead ratio from the provenance system to native.
2. We apply ID to the matrix. ID seeks to estimate a $m \times n$ matrix by retaining k of its columns and using a linear regression to estimate the remaining $n - k$ from those selected k columns. Cheng et al. [15] give a $O(k(m + n - k))$ algorithm for computing an optimal ID while keeping a reasonable⁵ L2 norm of the difference between the estimated and actual columns. Cheng’s procedure is implemented in `scipy.linalg.interpolative`⁶.
3. We select the k benchmarks corresponding to columns chosen by ID; these are the “best” k columns which minimize the error when predicting the $n - k$ columns under a specific metric.

- **Random search.** Random search proceeds like ID, but it selects k benchmarks randomly. Like ID, it computes a linear predictor for the $m - k$ benchmarks based on the k benchmarks. Then it evaluates the “goodness of fit” of that predictor, and repeats a fixed number of iterations, retaining the “best” k benchmark subset.

Cross-validation proceeds in the following manner, given m provenance systems, n benchmarks, and f features:

1. Separate 1 provenance system for testing from the $m - 1$ provenance systems used for training.
2. Use the $(m - 1) \times n$ training systems and $f \times n$ in one of the above algorithms to select the best $k < n$ benchmarks and compute predictors for the $n - k$ benchmarks.
3. Feed in the output of the $1 \times k$ testing provenance system on the selected k benchmarks into the algorithm, and let the algorithm estimate the $n - k$ unselected benchmarks.

⁵The best possible error for any rank- k factorization of a matrix is given by the $(k + 1)^{\text{th}}$ singular value. Since ID constrains the space of permissible factors, the L2 loss will be at least that. Cheng et al.’s ID method guarantees an error within a $\sqrt{1 + k(\min(m, n) - k)}$ factor of the $(k + 1)^{\text{th}}$ singular value. Read asymptotically, the bound asserts as the singular value decreases, so too does the L2 loss.

⁶See <https://docs.scipy.org/doc/scipy/reference/linalg.interpolative.html>

4. Score the difference between the algorithm’s prediction of the test system on the $n - k$ unselected benchmarks and its actual performance.
5. Repeat to 1 until all systems have been used for testing.

Note that during cross-validation, testing data (used in step 4) must not be allowed to “leak” into the training phase (step 2). For example, it would be invalid to do feature selection on all $m \times n$ data. Cross-validation is supposed to simulate the situation where one is testing on *truly novel* data.

We will evaluate these methods based on cross-validated root mean square-error (RMSE). Mean *square* error (MSE) preferable to mean *absolute* error (MAE) because it punishes outliers more. Imagine a benchmark suite which minimizes MAE and another which minimizes MSE. The MAE-minimizing suite one might be very accurate for most provenance systems, but egregiously wrong for a few; a MSE-minimizing suite may be “more wrong” on average, but the worst-case wouldn’t be as bad as the MAE one. As such, an MSE subset would be more practically useful for future publications to benchmark their new provenance systems.

We score the model on its ability to predict the logarithm of the ratio between a program running in provenance-generating and native systems. We use a ratio rather than the difference because some tests are very fast (10 seconds; so a 5 second predictive error matters a lot), while some tests are very slow (1000 seconds; so a 5 second predictive error matters very little). We predict the logarithm of the ratio, rather than the ratio directly because the ratio is multiplicative. Any real number is permissible; 0 indicates “nothing changed”, 1 indicates a speedup by a certain factor, and -1 indicates a slowdown *by the same factor*.

While cross-validation does punish model-complexity and overfitting to some extent, we will still take the number of parameters into account when deciding the “best” model in the interest of epistemic modesty. Preferring fewer parameters makes the model more generalizable on out-of-domain data, since even our full cross-validation data is necessarily incomplete.

3.5 Performance model

A related problem to subsetting is inferring a performance model. There are two motivations for inferring a performance model:

- A sysadmin may wish to provide a computational provenance capturing system to their institution, but getting approval to run new software on their system may be expensive (e.g., on highly secure systems, the sysadmin may need to acquire a security audit of the code before it can be approved for use). They may want to prospectively estimate the overhead of provenance collectors without having to install all the provenance collectors on their system, so they can select the optimal collector for their use-case.
- Inferring a provenance model may improve our understanding of the bottlenecks in provenance collectors.

A performance model should input features of a prospective workload and output the approximate overhead under different systems. A priori, provenance systems put a “tax” on certain syscalls (e.g., file I/O operations, process forks, process execs), because the system has to intercept and record these. Therefore, we expect a low-dimensional linear model (perhaps number of I/O operations

per second times a weight plus number of forks per second times another weight) would predict overhead optimally. To estimate this, we use the following models:

- **Ordinary least-squares (OLS) linear regression.** We estimate the runtime of each benchmark on each provenance system as a linear regression of the features of each benchmark, learning weights for each feature in each provenance system using ordinary least-squares. This would create a model like $\text{weight}_1 \cdot \text{feature}_1 + \text{weight}_2 \cdot \text{feature}_2 + \dots$. However, we can reduce its number of parameters, and thereby increase its out-of-domain generalizability, by the next two methods.
- **Low-rank linear regression.** To further reduce the number of parameters, we apply singular value decomposition (SVD) to create lossily-compressed representation of the learned weights. TODO: describe this model
- **OLS on a subset of features.** This method proceeds like the OLS regression, except it only uses a subset of the features, ignoring the rest. This is like doing a LASSO regression, but with multiple linear predictors sharing the same set of features (LASSO is usually described as solving for just one linear predictor). Unfortunately, we do not know an efficient algorithm like ID for selecting this subset. We tried to algorithms: greedy, which picks one additional feature that decreases loss the most until it has k features, and random, which selects a random k -sized subset.

We use as features:

- The number of $\$x$ -syscalls made per walltime second, where $\$x$ could be socket-related, file-related, reading-file-metadata, chmod, exec, clone, etc
- The number of syscalls per walltime second
- The amount of CPU time used per walltime second
- A constant fixed-cost per-execution

Like with benchmark minimization, we will use cross-validated RMSE errors in log of overhead ratio and the number of features in each model to select the best.

4 RESULTS

4.1 Selected provenance collectors

Table 2 shows the provenance collectors we collected and their qualitative features. We did not try reproducing provenance systems for the following reasons:

- **Not for Linux.** Our systems are Linux-based and Linux is used by many computational scientists. Therefore, we did not try to reproduce systems which were not Linux based.
- **VMs too slow.** Some provenance collectors require running the code in a virtual machine. We know a priori that these methods are prohibitively slow, with Panorama reporting 20x average overhead [57]. The provenance systems we are interested in have overheads in the 1.01x – 3x range.
- **Requires recompilation.** Some provenance collectors require the user to recompile their entire application and library stack. This is prohibitively onerous and negates the otherwise low cost of switching to system-level provenance we are pursuing.

- **Requires special hardware.** Some methods require certain CPUs, e.g., Intel CPUs for a dynamic instrumentation tool called Intel PIN. Being limited to certain CPUs violates our goal of promulgating reproducibility to as many people as possible.
- **No source.** We searched the original papers, GitHub, BitBucket, Google, and emailed the first author (CCing the others). If we still could not find the source code for a particular provenance collector, we cannot reproduce it. Note, however, that RecProv is implemented using rr, so we can use rr as a lower-bound for RecProv.
- **Ancient kernel (Hi-Fi and LPM/ProvMon).** Some provenance systems are implemented as patches into the Linux kernel. As time passed, these grew out-of-date with modern Linux kernels. We deprioritized the implementation of these methods because they require extremely old kernels, and thus may not be worthy of use in practical systems. Even conservative Linux distributions like CentOS 7 use Linux 3.10. This difficulty is not a bug in our study, but reflects an underlying reality that modified kernels are less likely to be maintained. However, if we had more time, we would want to reproduce these systems too.
- **Need more time for kernel (CamFlow).** Provenance systems that are implemented as patches on the Linux kernel are difficult to deploy, even if the kernel is recent. We could implement experimental infrastructure to run the modified kernel on “bare metal”, which is difficult to set up and dangerous, or we could run the modified kernel in a traditional virtual machine, which would distort the runtimes non-linearly from native performance and invalidate the predictive performance model, or a cycle-accurate virtual machine, which would prohibitively slow down the benchmark suite to the point that we would have much less data. Nevertheless, we would like to implement this system in a cycle-accurate virtual machine like gem5, if we had more time.
- **Not reproducible (OPUS).** We tried to get this provenance system to run with several weeks of effort, we emailed the original authors and other authors who used this system, and we left a GitHub issue describing the expected and actual results⁷, however we still could not get the system to run properly. If we had more time, perhaps we could debug the issue.
- **Needs more time (DTrace, SPADE, eBPF/bpftrace).** We simply needed more time to implement these provenance collectors.
- **Reproduced/excluded (ltrace).** ltrace is an off-the-shelf tool, available in most Linux package repositories. However, we found it could not handle several of our benchmark workloads. We localized the problem to the following code⁸:

```
/* FIXME: not good -- should use dynamic allocation. 19990703 mortene. */
if (proc->callstack_depth == MAX_CALLDEPTH - 1) {
    fprintf(stderr, "%s: Error: call nesting too deep!\n", __func__);
    abort();
    return;
}
```

⁷See <https://github.com/dtg-FRESCO/opus/issues/1>

⁸See https://gitlab.com/cespedes/ltrace/-/blob/8eabf684ba6b11ae7a1a843aca3c0657c6329d73/handle_event.c#L775

Table 2: Provenance collectors mentioned in primary and secondary studies in our search results.

Tool	Collection method	Status
strace	Tracing	Reproduced
fsatrace	Library ins.	Reproduced
ReproZip [16]	Tracing	Reproduced
Sciunit2 [53]	Tracing	Reproduced
rr [42]	Tracing	Reproduced
CDE [24]	Tracing	Reproduced
ltrace	Tracing	Reproduced/excluded
SPADE [21]	Audit, FS, or compile-time	Needs more time
DTrace [1]	Audit	Needs more time
eBPF/bpftrace	Audit	Needs more time
CamFlow [43]	Kernel ins.	Needs more time for kern.
OPUS [8]	Library ins.	Not reproducible
Hi-Fi [45]	Kernel ins.	Ancient kernel (3.2.0)
LPM/ProvMon [10]	Kernel ins.	Ancient kernel (2.6.32)
RecProv [27]	Tracing	No source
LPROV [56]	Kernel mod., lib. ins.	No source
S2Logger [51]	Kernel mod.	No source
ProTracer [37]	Kernel mod.	No source
FIPS [52]	FS ins.	No source
PANDDE [19]	FS ins.	No source
PASS/Pasta [41]	Kernel ins., FS ins.	No source
PASSv2/Lasagna [40]	Kernel ins., FS, lib. ins.	No source
Lineage FS [47]	Kernel ins.	No source
RTAG [26]	Kernel ins.	No source
BEEP [34]	Dyn./static binary ins.	Requires HW
libdft [29]	Dyn. binary ins.	Requires HW
RAIN [25]	Dyn. bin., kernel, lib. ins.	Requires HW
DataTracker [50]	Dinamic binary ins.	Requires HW
MPI [36]	Compile-time ins.	Requires recompilation
LDX [31]	Compile-time ins.	Requires recompilation
Panorama [57]	VM ins.	VMs are too slow
PROV-Tracer [49]	VM ins.	VMs are too slow
ETW [5]	Audit	Not for Linux
Sysmon [39]	Audit	Not for Linux
TREC [55]	Audit	Not for Linux
URSprung [46]	Audit, file-system ins.	Not for Linux ⁹
Ma et al. [35]	Kernel ins.	Not for Linux

4.2 Implemented benchmarks

Of these, table 3 shows the benchmarks used to evaluate each tool, of which there are quite a few. First, we eliminated several benchmarks from this set as non-starters for the reasons described in ?? . Then, we prioritized implementing frequently-used benchmarks, easy-to-implement benchmarks, and benchmarks which we believe have value in representing a computational science use-case.

- **HTTP/FTP servers/clients/traffic.** The most common benchmark class from prior work, HTTP servers/traffic, HTTP servers/clients, FTP servers/traffic, and FTP servers/clients are popular because prior work focuses overwhelmingly on provenance for the sake of security (auditing, intrusion detection, or digital forensics). Therefore, we implement HTTP servers with simulated traffic, HTTP servers with specific clients, and FTP servers with simulated traffic. Although we wanted to implement FTP servers with specific clients, all of those executions reach the callstack depth in ltrace, even when fetching a single file, so we had to drop those. We do not think they are substantially different from the FTP servers with simulated traffic.
- **Compiling packages.** Compiling packages from source is a common operation in computational science, therefore

⁹URSprung depends on IBM Spectrum Scale to get directory change notifications, so it is not for a generic Linux system.

¹⁰LogGC measures the offline running time and size of garbage collected logs; there is no comparison to native would be applicable.

Table 3: Benchmarks used in various provenance publications.

Publication	Benchmarks	Comparisons
TREC [55]	open/close, compile Apache, compile LaTeX doc	Native
PASS [41]	BLAST	Native ext2
Panorama [57]	curl, scp, gzip, bzip2	Native
PASSv2 [40]	BLAST, compile Linux, Postmark, Mercurial, Kepler	Native ext3, NFS
SPADEv2 [21]	BLAST, compile Apache, Apache	Native
Hi-Fi [45]	lmbench, compile Linux, Postmark	Native
libdft [29]	scp, tar, gzip, bzip2 x extract, compress	PIN
LogGC [33]	RUBiS, Firefox, MC, Pidgin, Pine, Proftpd, Sendmail, sshd, vim, w3m, wget, xpdf, yafc, Audacious, bash, Apache, mysqld	None ¹⁰
LPM/ProvMon [10]	lmbench, compile Linux, Postmark, BLAST	Native
Ma et al. [35]	TextTransfer, Chromium, DrawTool, NetFTP, AdvancedFTP, Apache, IE, Paint, Notepad, Notepad++, simplehttp, Sublime Text	Native
ProTracer [37]	Apache, miniHTTP, ProFTPD, Vim, Firefox, w3m, wget, mplayer, Pine, xpdf, MC, yafc	Auditd, BEEP
LDX [31]	SPEC CPU 2006, Firefox, lynx, nginx, tnftp, sysstat, gif2png, mp3info, prozilla, yopweb, ngired, gocr, Apache, pbzip2, pigz, axel, x264	Native
PANDDE [19]	ls, cp, cd, lpr	Native
MPI [36]	Apache, bash, Evince, Firefox, Krusader, wget, most, MC, mplayer, MPV, nano, Pine, ProFTPD, SKOD, Tiny-HTTPD, Transmission, Vim, w3m, xpdf, Yafc	Audit, LPM-HiFi
CamFlow [43]	lmbench, postmark, unpack kernel, compile Linux, Apache, Memcache, redis, php, pybench	Native
BEEP [34]	Apache, Vim, Firefox, wget, Cherokee, w3m, ProFTPD, yafc, Transmission, Pine, bash, mc, sshd, sendmail	Native
RAIN [25]	SPEC CPU 2006, cp linux, wget, compile libc, Firefox, SPLASH-3	Native
Sciunit [53]	VIC, FIE	Native
LPROV [56]	Apache, simplehttp, proftpd, sshd, firefox, filezilla, lynx, links, w3m, wget, ssh, pine, vim, emacs, xpdf	Native
MCI [32]	Firefox, Apache, Lighttpd, nginx, ProFTPD, CUPS, vim, elinks, alpine, zip, transmission, lftp, yafc, wget, ping, procps	BEEP
RTAG [26]	SPEC CPU 2006, scp, wget, compile llvm, Apache	RAIN
URSPRING [46]	open/close, fork/exec/exit, pipe/dup/close, socket/connect, CleanML, Vanderbilt, Spark, ImageML	Native, SPADE

we implemented as many of these as we could and implemented some of our own. However, compiling LLVM takes more than twice as long as the longest benchmark, so we excluded LLVM specifically from the benchmark suite. We implemented a pattern for compiling packages from Spack which discounts the time taken to download sources, counting only the time taken to unpack, patch, configure, compile, link, and install them. We try compiling Python, Boost, HDF5, glibc, Apache HTTPd, and Perl. We would like to implement compilation for more packages, in particular xSDK [9] packages.

- **Browsers.** Implementing headless for browsers in “batch-mode” without GUI interaction is not impossibly difficult, but non-trivial. Furthermore, we deprioritized this benchmark because few computational science applications resemble the workload of a web browser. If we had more time, we would implement this benchmark too.
- **Un/archive.** Archive and unarchiving would be a common task for retrieving data or source code. We benchmark un/archiving several archives with several compression algorithms. Choosing a compression algorithm may turn an otherwise I/O-bound workload to a CPU-bound workload, which would make the impact of provenance tracing smaller.

- **I/O microbenchmarks (lmbench, postmark, custom).** These could be informative for explicating which I/O operations are most affected. Prior work uses lmbench [?], which focuses on syscalls generally, Postmark [28], which focuses on I/O operations, and custom benchmarks, for example running open/close in a tight loop. We use the specific lmbench cases from prior work, which is mostly the latency benchmarks with a few bandwidth benchmarks. Most provenance systems do not affect the bandwidth; it doesn't matter *how much* this process writes to that file, just *that* this process wrote to that file.
- **BLAST.** BLAST [7] is a search for a fuzzy string in a protein database. Many prior works use this as a file-read heavy benchmark, as do we. This code in particular resembles a computational science workload. However, unlike prior work, we split the benchmark into each of each subtasks; provenance may have a greater overhead on certain subtasks.
- **CPU benchmarks.** SPEC CPU INT 2006 [?] and SPLASH-3 [?] test CPU performance. Provenance collection would not slow down CPU integer processing at all, since it only intercepts I/O related library- or system-calls. Therefore, we deprioritized these benchmarks, and we did not have enough time to write reproducible build scripts for SPEC CPU INT 2006.
- **Sendmail.** Sendmail is a quite old mail server program. Mail servers do not resemble a computational science workload, and it is unclear what workload we would run against the server. Therefore, we deprioritized this benchmark and did not have time to implement it.
- **VCS checkouts.** VCS checkouts are a common computational science operation. Prior work uses Mercurial, but we implemented Mercurial and Git. We simply clone a repository (untimed) and run \$vcs checkout for random commits in the repository (timed).
- **Data processing/machine-learning Workflows.** VIC, FIE, ImageML, and Spark are real-world examples of scientific workflows. We would like to implement these, but reproducing those workflows is non-trivial; they each require their own computational stack. For FIE, in particular, there is no script that glues all of the operations together; we would have to read the publication [12] which FIE supports to understand the workflow, and write our own script which glues the operations together.
- **Coreutils and other utilities (bash, cp, ls, procps).** We do not see a huge representative value in these benchmarks that would not already be gleaned from lmbench, but due to its simplicity, we implemented it anyway. For bash, we do not know what workload prior works are using, but we test the speed of incrementing an integer and changing directories (cd).
- The rest of the programs are mostly specific desktop applications used only in one prior work. These would likely not yield any insights not already yielded by the benchmarks we implemented, and for each one we would need to build it from source, find a workload for it, and take the time to run it. They weigh little in the argument that our benchmark suite represents prior work, since they are only used in one

Table 4: Benchmarks implemented by this work.

Prior works	This work	Benchmark group and examples from prior work
10	yes (5/7 servers)	HTTP server/traffic (Apache httpd, miniHTTP, simplehttp, lighttpd, Nginx, tinyhttpd, cherokee x apachebench)
9	yes (3/4 clients)	HTTP server/client (simplehttp x curl, wget, prozilla, axel)
8	yes (3/5 orig + 4 others)	Compile user packages (Apache, LLVM, glibc, Linux, LaTeX document)
5	yes	FTP server/traffic (ProFTPD x ftpbench)
8	no	Browsers (Firefox, Chromium x Sunspider)
5	yes	Un/archive (compress, decompress x tar x nothing, bzip2, pbzip, gzip, pigz, zip)
5	yes	I/O microbenchmarks (Postmark, lmbench, custom)
4	yes	BLAST
3	no	CPU benchmarks (SPEC CPU INT 2006, SPLASH-3)
3	yes	Coreutils and other utils (bash, cp, ls, procps)
2	no	Sendmail
1	yes	VCS checkouts (Mercurial)
1	no	Machine learning workflows (CleanML, Spark)
1	no	Data processing workflows (VIC, FIE)
1	no	RUBiS
1	no	x264
1	no	mysqld
1	no	gocr
1	no	Memcache
1	no	Redis
1	no	php
1	no	pybench
1	no	ImageML
1	no	ping
1	no	mp3info
1	no	ngircd
1	no	CUPS

prior work. We would like to implement these, but they have a lower priority.

4.3 Subsetted benchmarks

4.4 Quantitative performance comparison

4.5 Predictive model

5 ANALYSIS

5.1 Threats to validity

6 FUTURE WORK

7 CONCLUSION

8 REFERENCES

REFERENCES

- [1] [n. d.]. About DTrace.
- [2] [n. d.]. BPF Documentation. <https://docs.kernel.org/bpf/index.html>.
- [3] [n. d.]. FUSE. <https://www.kernel.org/doc/html/latest/filesystems/fuse.html>.
- [4] [n. d.]. Ptrace. <https://man7.org/linux/man-pages/man2/ptrace.2.html>.
- [5] 2021. Event Tracing - Win32 Apps. <https://learn.microsoft.com/en-us/windows/win32/etw/event-tracing-portal>.
- [6] ACM Inc. staff. 2020. Artifact Review and Badging. <https://www.acm.org/publications/policies/artifact-review-and-badging-current>.
- [7] Stephen F. Altschul, Warren Gish, Webb Miller, Eugene W. Myers, and David J. Lipman. 1990. Basic Local Alignment Search Tool. *Journal of Molecular Biology* 215, 3 (Oct. 1990), 403–410. [https://doi.org/10.1016/S0022-2836\(05\)80360-2](https://doi.org/10.1016/S0022-2836(05)80360-2)
- [8] Nikilesh Balakrishnan, Thomas Bytheway, Ripduman Sohan, and Andy Hopper. 2013. {OPUS}: A Lightweight System for Observational Provenance in User Space. In *5th USENIX Workshop on the Theory and Practice of Provenance (TaPP 13)*.
- [9] Roscoe Bartlett, Irina Demeshko, Todd Gamblin, Glenn Hammond, Michael Allen Heroux, Jeffrey Johnson, Alicia Klinvex, Xiaoye Li, Lois Curfman McInnes, J. David Moulton, Daniel Osei-Kuffuor, Jason Sarich, Barry Smith, James Willenbring, and Ulrike Meier Yang. 2017. xSDK Foundations: Toward an Extreme-scale Scientific Software Development Kit. *Supercomputing Frontiers and Innovations* 4, 1 (Feb. 2017), 69–82. <https://doi.org/10.14529/jsfi170104>

- [10] Adam Bates, Dave (Jing) Tian, Kevin R. B. Butler, and Thomas Moyer. 2015. Trustworthy {Whole-System} Provenance for the Linux Kernel. In *24th USENIX Security Symposium (USENIX Security 15)*. 319–334.
- [11] Dirk Beyer, Stefan Löwe, and Philipp Wendler. 2019. Reliable Benchmarking: Requirements and Solutions. *Int J Softw Tools Technol Transfer* 21, 1 (Feb. 2019), 1–29. <https://doi.org/10.1007/s10009-017-0469-y>
- [12] Mirza M. Billah, Jonathan L. Goodall, Ujjwal Narayan, Bakinam T. Essawy, Venkat Lakshmi, Arcot Rajasekar, and Reagan W. Moore. 2016. Using a Data Grid to Automate Data Preparation Pipelines Required for Regional-Scale Hydrologic Modeling. *Environmental Modelling & Software* 78 (April 2016), 31–39. <https://doi.org/10.1016/j.envsoft.2015.12.010>
- [13] Bruno Cartaxo, Gustavo Pinto, and Sergio Soares. 2018. The Role of Rapid Reviews in Supporting Decision-Making in Software Engineering Practice. In *Proceedings of the 22nd International Conference on Evaluation and Assessment in Software Engineering 2018 (EASE '18)*. Association for Computing Machinery, New York, NY, USA, 24–34. <https://doi.org/10.1145/3210459.3210462>
- [14] Bruno Cartaxo, Gustavo Pinto, and Sergio Soares. 2020. Rapid Reviews in Software Engineering. In *Contemporary Empirical Methods in Software Engineering*, Michael Felderer and Guilherme Horta Travassos (Eds.). Springer International Publishing, Cham, 357–384. https://doi.org/10.1007/978-3-030-32489-6_13
- [15] H. Cheng, Z. Gimbutas, P. G. Martinsson, and V. Rokhlin. 2005. On the Compression of Low Rank Matrices. *SIAM J. Sci. Comput.* 26, 4 (Jan. 2005), 1389–1404. <https://doi.org/10.1137/030602678>
- [16] Fernando Chirigati, Rémi Rampin, Dennis Shasha, and Juliana Freire. 2016. ReproZip: Computational Reproducibility With Ease. In *Proceedings of the 2016 International Conference on Management of Data (SIGMOD '16)*. Association for Computing Machinery, New York, NY, USA, 2085–2088. <https://doi.org/10.1145/2882903.2899401>
- [17] Christian Collberg and Todd A. Proebsting. 2016. Repeatability in Computer Systems Research. *Commun. ACM* 59, 3 (Feb. 2016), 62–69. <https://doi.org/10.1145/2812803>
- [18] Fred D. Davis. 1985. *A Technology Acceptance Model for Empirically Testing New End-User Information Systems: Theory and Results*. Thesis. Massachusetts Institute of Technology.
- [19] Daren Fadolalkarim, Asmaa Sallam, and Elisa Bertino. 2016. PANDDE: Provenance-based Anomaly Detection of Data Exfiltration. In *Proceedings of the Sixth ACM Conference on Data and Application Security and Privacy (CO-DASPY '16)*. Association for Computing Machinery, New Orleans Louisiana USA, 267–276. <https://doi.org/10.1145/2857705.2857710>
- [20] Juliana Freire, David Koop, Emanuele Santos, and CI Silva. 2008. Provenance for Computational Tasks: A Survey. *Comput. Sci. Eng.* 10, 3 (May 2008), 11–21. <https://doi.org/10.1109/MCSE.2008.79>
- [21] Ashish Gehani and Dawood Tariq. 2012. SPADE: Support for Provenance Auditing in Distributed Environments. In *Middleware 2012 (Lecture Notes in Computer Science)*, Priya Narasimhan and Peter Triantafillou (Eds.). Springer, Berlin, Heidelberg, 101–120. https://doi.org/10.1007/978-3-642-35170-9_6
- [22] Gooch. [n.d.]. Overview of the Linux Virtual File System. <https://docs.kernel.org/filesystems/vfs.html>
- [23] Samuel Grayson, Darko Marinov, Daniel S. Katz, and Reed Milewicz. 2023. Automatic Reproduction of Workflows in the Snakemake Workflow Catalog and NF-Core Registries. In *Proceedings of the 2023 ACM Conference on Reproducibility and Replicability (ACM REP '23)*. Association for Computing Machinery, New York, NY, USA, 74–84. <https://doi.org/10.1145/3589806.3600037>
- [24] Philip Guo and Dawson Engler. 2011. CDE: Using System Call Interposition to Automatically Create Portable Software Packages. In *2011 USENIX Annual Technical Conference*. USENIX, Portland, OR, USA.
- [25] Yang Ji, Sangho Lee, Evan Downing, Weiren Wang, Mattia Fazzini, Taesoo Kim, Alessandro Orso, and Wenke Lee. 2017. RAIN: Refinable Attack Investigation with On-demand Inter-Process Information Flow Tracking. In *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security (CCS '17)*. Association for Computing Machinery, New York, NY, USA, 377–390. <https://doi.org/10.1145/3133956.3134045>
- [26] Yang Ji, Sangho Lee, Mattia Fazzini, Joey Allen, Evan Downing, Taesoo Kim, Alessandro Orso, and Wenke Lee. 2018. Enabling Refinable {Cross-Host} Attack Investigation with Efficient Data Flow Tagging and Tracking. In *27th USENIX Security Symposium (USENIX Security 18)*. 1705–1722.
- [27] Yang Ji, Sangho Lee, and Wenke Lee. 2016. RecProv: Towards Provenance-Aware User Space Record and Replay. In *Provenance and Annotation of Data and Processes (Lecture Notes in Computer Science)*, Marta Mattoso and Boris Glavic (Eds.). Springer International Publishing, Cham, 3–15. https://doi.org/10.1007/978-3-319-40593-3_1
- [28] Jeffrey Katcher. 2005. *PostMark: A New File System Benchmark*. Technical Report TR3022.
- [29] Vasileios P. Kemerlis, Georgios Portokalidis, Kangkook Jee, and Angelos D. Keromytis. 2012. Libdft: Practical Dynamic Data Flow Tracking for Commodity Systems. In *Proceedings of the 8th ACM SIGPLAN/SIGOPS Conference on Virtual Execution Environments (VEE '12)*. Association for Computing Machinery, New York, NY, USA, 121–132. <https://doi.org/10.1145/2151024.2151042>
- [30] Jim Keniston, Prasanna S Panchamukhi, and Masami Hiramatsu. [n. d.]. Kernel Probes (Kprobes). <https://www.kernel.org/doc/html/latest/trace/kprobes.html>.
- [31] Yonghui Kwon, Dohyeon Kim, William Nick Sumner, Kyungtae Kim, Brendan Saltaformaggio, Xiangyu Zhang, and Dongyan Xu. 2016. LDX: Causality Inference by Lightweight Dual Execution. In *Proceedings of the Twenty-First International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS '16)*. Association for Computing Machinery, New York, NY, USA, 503–515. <https://doi.org/10.1145/2872362.2872395>
- [32] Yonghui Kwon, Fei Wang, Weihang Wang, Kyu Hyung Lee, Wen-Chuan Lee, Shiqing Ma, Xiangyu Zhang, Dongyan Xu, Somesh Jha, Gabriela Ciocarlie, Ashish Gehani, and Vinod Yegneswaran. 2018. MCI: Modeling-based Causality Inference in Audit Logging for Attack Investigation. In *Proceedings 2018 Network and Distributed System Security Symposium*. Internet Society, San Diego, CA. <https://doi.org/10.14722/ndss.2018.23306>
- [33] Kyu Hyung Lee, Xiangyu Zhang, and Dongyan Xu. 2013. LogGC: Garbage Collecting Audit Log. In *Proceedings of the 2013 ACM SIGSAC Conference on Computer & Communications Security (CCS '13)*. Association for Computing Machinery, New York, NY, USA, 1005–1016. <https://doi.org/10.1145/2508859.2516731>
- [34] Kyu Hyung Lee, Xiangyu Zhang, and Dongyan Xu. 2017. High Accuracy Attack Provenance via Binary-based Execution Partition. In *Proceedings of the 2017 Network and Distributed System Security (NDSS) Symposium*.
- [35] Shiqing Ma, Kyu Hyung Lee, Chung Hwan Kim, Junghwan Rhee, Xiangyu Zhang, and Dongyan Xu. 2015. Accurate, Low Cost and Instrumentation-Free Security Audit Logging for Windows. In *Proceedings of the 31st Annual Computer Security Applications Conference (ACSAC '15)*. Association for Computing Machinery, New York, NY, USA, 401–410. <https://doi.org/10.1145/2818000.2818039>
- [36] Shiqing Ma, Juan Zhai, Fei Wang, Kyu Hyung Lee, Xiangyu Zhang, and Dongyan Xu. 2017. [MPI]: Multiple Perspective Attack Investigation with Semantic Aware Execution Partitioning. In *26th USENIX Security Symposium (USENIX Security 17)*. 1111–1128.
- [37] Shiqing Ma, Xiangyu Zhang, and Dongyan Xu. 2016. ProTracer: Towards Practical Provenance Tracing by Alternating Between Logging and Tainting. In *Proceedings 2016 Network and Distributed System Security Symposium*. Internet Society, San Diego, CA. <https://doi.org/10.14722/ndss.2016.23350>
- [38] Ashish Bharadwaj Madabhushana. 2021. Configure Linux System Auditing with Auditd.
- [39] markruss. 2023. Sysmon - Sysinternals. <https://learn.microsoft.com/en-us/sysinternals/downloads/sysmon>.
- [40] Kiran-Kumar Muniswamy-Reddy, Uri Braun, David A. Holland, Peter Macko, Diana Maclean, Daniel Margo, Margo Seltzer, and Robin Smogor. 2009. Layering in Provenance Systems. In *Proceedings of the 2009 Conference on USENIX Annual Technical Conference (USENIX'09)*. USENIX Association, USA, 10. <https://doi.org/10.5555/1855807.1855817>
- [41] Kiran-Kumar Muniswamy-Reddy, David A Holland, Uri Braun, and Margo Seltzer. 2006. Provenance-Aware Storage Systems. In *2006 USENIX Annual Technical Conference*.
- [42] Robert O'Callahan, Chris Jones, Nathan Froyd, Kyle Huey, Albert Noll, and Nimrod Partush. 2017. Engineering Record And Replay For Deployability: Extended Technical Report. <https://doi.org/10.48550/arXiv.1705.05937> arXiv:1705.05937 [cs]
- [43] Thomas Pasquier, Xueyuan Han, Mark Goldstein, Thomas Moyer, David Eysers, Margo Seltzer, and Jean Bacon. 2017. Practical Whole-System Provenance Capture. In *Proceedings of the 2017 Symposium on Cloud Computing (SoCC '17)*. Association for Computing Machinery, New York, NY, USA, 405–418. <https://doi.org/10.1145/3127479.3129249>
- [44] João Felipe Pimentel, Juliana Freire, Leonardo Murta, and Vanessa Braganholo. 2019. A Survey on Collecting, Managing, and Analyzing Provenance from Scripts. *ACM Comput. Surv.* 52, 3 (June 2019), 47:1–47:38. <https://doi.org/10.1145/3311955>
- [45] Devin J. Pohly, Stephen McLaughlin, Patrick McDaniel, and Kevin Butler. 2012. Hi-Fi: Collecting High-Fidelity Whole-System Provenance. In *Proceedings of the 28th Annual Computer Security Applications Conference (ACSAC '12)*. Association for Computing Machinery, New York, NY, USA, 259–268. <https://doi.org/10.1145/2420950.2420989>
- [46] Lukas Rupperecht, James C. Davis, Constantine Arnold, Yaniv Gur, and Deepavali Bhagwat. 2020. Improving Reproducibility of Data Science Pipelines through Transparent Provenance Capture. *Proc. VLDB Endow.* 13, 12 (Aug. 2020), 3354–3368. <https://doi.org/10.14778/3415478.3415556>
- [47] Can Sar and Pei Cao. [n. d.]. Lineage File System. [n. d.].
- [48] Holger J. Schünemann and Lorenzo Moja. 2015. Reviews: Rapid! Rapid! Rapid! ... and Systematic. *Systematic Reviews* 4, 1 (Jan. 2015), 4. <https://doi.org/10.1186/2046-4053-4-4>
- [49] Manolis Stamatogiannakis, Paul Groth, and Herbert Bos. 2015. Decoupling Provenance Capture and Analysis from Execution. In *Proceedings of the 7th USENIX Conference on Theory and Practice of Provenance (TaPP'15)*. USENIX Association, USA, 3.
- [50] Manolis Stamatogiannakis, Paul Groth, and Herbert Bos. 2015. Looking Inside the Black-Box: Capturing Data Provenance Using Dynamic Instrumentation. In

- Provenance and Annotation of Data and Processes (Lecture Notes in Computer Science), Bertram Ludäscher and Beth Plale (Eds.). Springer International Publishing, Cham, 155–167. https://doi.org/10.1007/978-3-319-16462-5_12
- [51] Chun Hui Suen, Ryan K.L. Ko, Yu Shyang Tan, Peter Jagadpramana, and Bu Sung Lee. 2013. S2Logger: End-to-End Data Tracking Mechanism for Cloud Data Provenance. In *2013 12th IEEE International Conference on Trust, Security and Privacy in Computing and Communications*. 594–602. <https://doi.org/10.1109/TrustCom.2013.73>
- [52] Salmin Sultana and Elisa Bertino. 2013. A File Provenance System. In *Proceedings of the Third ACM Conference on Data and Application Security and Privacy (CO-DASPY '13)*. Association for Computing Machinery, New York, NY, USA, 153–156. <https://doi.org/10.1145/2435349.2435368>
- [53] Dai Hai Ton That, Gabriel Fils, Zhihao Yuan, and Tanu Malik. 2017. Sciunits: Reusable Research Objects. In *2017 IEEE 13th International Conference on E-Science (e-Science)*. 374–383. <https://doi.org/10.1109/eScience.2017.51>
- [54] Ana Trisovic, Matthew K. Lau, Thomas Pasquier, and Mercè Crosas. 2022. A Large-Scale Study on Research Code Quality and Execution. *Sci Data* 9, 1 (Feb. 2022), 60. <https://doi.org/10.1038/s41597-022-01143-6>
- [55] Amin Vahdat and Thomas Anderson. 1998. Transparent Result Caching. In *Proceedings of the Annual Conference on USENIX Annual Technical Conference (ATEC '98)*. USENIX Association, USA, 3.
- [56] Fei Wang, Yonghui Kwon, Shiqing Ma, Xiangyu Zhang, and Dongyan Xu. 2018. Lprov: Practical Library-aware Provenance Tracing. In *Proceedings of the 34th Annual Computer Security Applications Conference (ACSAC '18)*. Association for Computing Machinery, New York, NY, USA, 605–617. <https://doi.org/10.1145/3274694.3274751>
- [57] Heng Yin, Dawn Song, Manuel Egele, Christopher Kruegel, and Engin Kirda. 2007. Panorama: Capturing System-Wide Information Flow for Malware Detection and Analysis. In *Proceedings of the 14th ACM Conference on Computer and Communications Security (CCS '07)*. Association for Computing Machinery, New York, NY, USA, 116–127. <https://doi.org/10.1145/1315245.1315261>
- [58] Jun Zhao, Jose-Manuel Gomez-Perez, Khalid Belhajjame, Graham Klyne, Esteban Garcia-cuesta, Aleix Garrido, Kristina Hettne, Marco Roos, David De Roure, and Carole Goble. 2012. Why Workflows Break — Understanding and Combating Decay in Taverna Workflows. In *2012 IEEE 8th International Conference on E-Science (e-Science)*. IEEE, Chicago, IL, 9. <https://doi.org/10.1109/eScience.2012.6404482>