

System-level provenance tracers

by Samuel Grayson, Reed Milewicz, Daniel S. Katz, and
Darko Marinov

Takeaways

1. Provenance is useful for computational scientists
2. Current state-of-the-art not practical use
3. But there is hope

What is provenance?

A Vase of Flowers / Pieces

[Reports](#) [Edit](#) [Copy](#) [Register Sale](#) [Donate](#)



Artist: [Margareta Haverman](#)
Size (h x w): 31.25 x 23.75 in
Medium: Oil on Wood
Subject Matter: Still Life
Type: Painting
Value: \$25,000.00
Insurance Value: \$25,000.00
Creation Date: 1716

Purchase Info

Purchase Date: 1917
Purchase Price: \$5,000.00
Purchase Location: [edit](#)
Seller: [John Johnston](#)

Current Location

[Boston Office](#) / Northeast hallway
[Jun 11, 2018](#) [\(edit\)](#)

[Assign To New Location](#)

Appraisal History [Add Appraisal](#)

Appraiser	Value	Date	
Katerina Jenkins	\$25,000.00	February 12, 2018	View

Location History [New Location Record](#)

Location	Dates	Current	
Boston Office	Jun 11, 2018 -	✓	Edit
Corporate Headquarters	Oct 01, 2017 - Nov 04, 2014 almost 3 years		Edit
The Met Fifth Avenue	Nov 04, 2014 - Jun 19, 2015 8 months		Edit

Limited Edition Runs [New Run](#)

Additional Files [Add File](#)



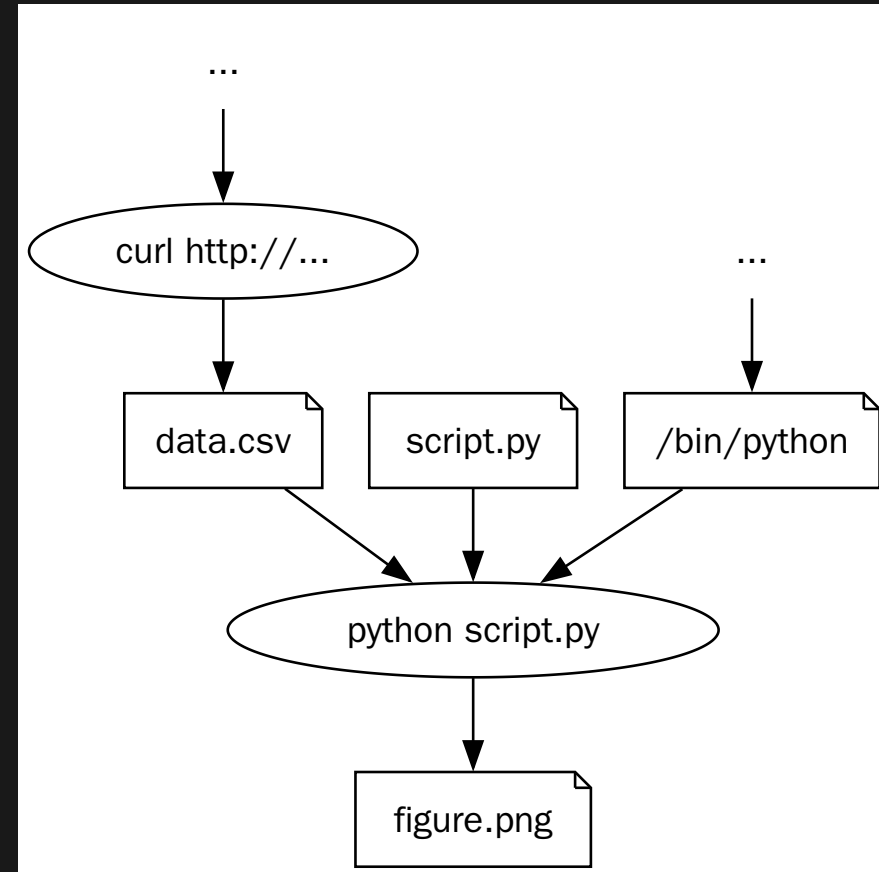
What is computational provenance?

1. Process by which a comp. artifact (often a file) was generated
2. The inputs to that process
3. The provenance of those inputs (recursively)

Expressible as a graph whose nodes are processes or files

Comp. Provenance Example

```
curl http://...  
python script.py
```



What is it good for?

- Record/replay reproducibility (*)
- What parameter of X did this output use? Digital notebook
- How does the ancestry of these two objects differ?
- Create a Spack/Nix/Guix/Makeflow automatically (*)

What is the state of the art?

How to collect provenance data?


- Language-level (rdtLite)
- Switch to workflows
- System-level (*)

How to collect provenance data?

- Language-level (rdtLite)
- Switch to workflows
- System-level (*)

Rapid review

Search queries: "Computational provenance" &
"System-level provenance"



Feasibility study

Linux sys-level prov tracers

Provenance tracer	Strategy	Modified kernel?	Privileges required?
PASSv2, CamFlow, ...	Modify kernel	yes	yes
SPADE, eBPF, ...	Built-in auditing	no	yes
Sciunit, ReproZip, ...	User tracing	no	no

Linux sys-level prov tracers

Provenance tracer	Strategy	Modified kernel?	Privileges required?
PASSv2, CamFlow, ...	Modify kernel	yes	yes
SPADE, eBPF, ...	Built-in auditing	no	yes
Sciunit, ReproZip, ...	User tracing	no	no

Performance analysis

Selected provenance tracers

- RecProv: Could not locate source code or authors
- OPUS: Could not reproduce at all (Py 2.7)
- CDE: Sometimes crashes
- PTU/Sciunit: Sometimes out-of-mem, false negatives
- ReproZip: Works
- CARE: Works

Related tools

- Strace: ptrace syscall tracking
- RR-debugger: ptrace record-replay
- fsatrace: LD_PRELOAD file tracking

Benchmarks used by prior works

Prior pubs.	This work	Benchmark
12	Yes	HTTP servers/clients
9	No	Web browsers
6	Yes	FTP servers/clients
5	Yes	Un/archive
5	Yes	BLAST
4	Yes	Postmark
3	Yes	Imbench

Performance

Benchmark	Native	fsatrace	CARE	strace	RR	ReproZip
	None	Lib. interp.	Ptrace	Ptrace	Ptrace	Ptrace
BLAST	0	0	2	2	93	8
Tar Unarchive	0	4	44	114	195	149
Python import	0	5	85	84	150	346
VCS checkout	0	5	71	160	177	428
Compile w/Spack	0	-1	119	111	562	359
Postmark	0	2	231	650	259	1733
cp	0	37	641	380	232	5791
Others not shown
Geometric mean	0	0	45	66	146	193

Discussion

Security

Un-circumventable at all costs

Assume we have root on baremetal

Use servers (like HTTP) as benchmarks

Comp. sci.

Can assume the codes are not intentionally deceptive

Don't have root on large-scale computers

Use workflows and comp experiments as benchmarks

Why aren't prov tracers used

- Require root
- Fragile/crashy
- Hard to install
- Too slow
- Happy with existing system

But there is hope!

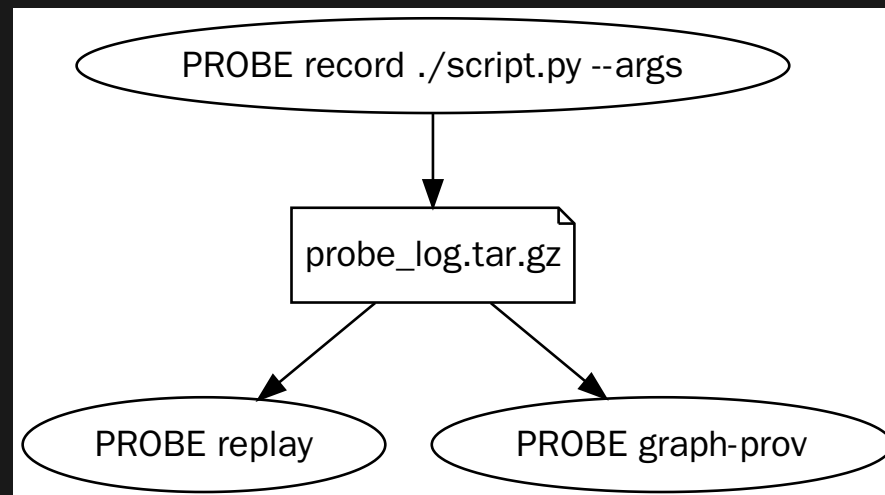
- Computational is niche
- Lib. interposition understudied in prior work
- Common set of benchmarks

Conclusion

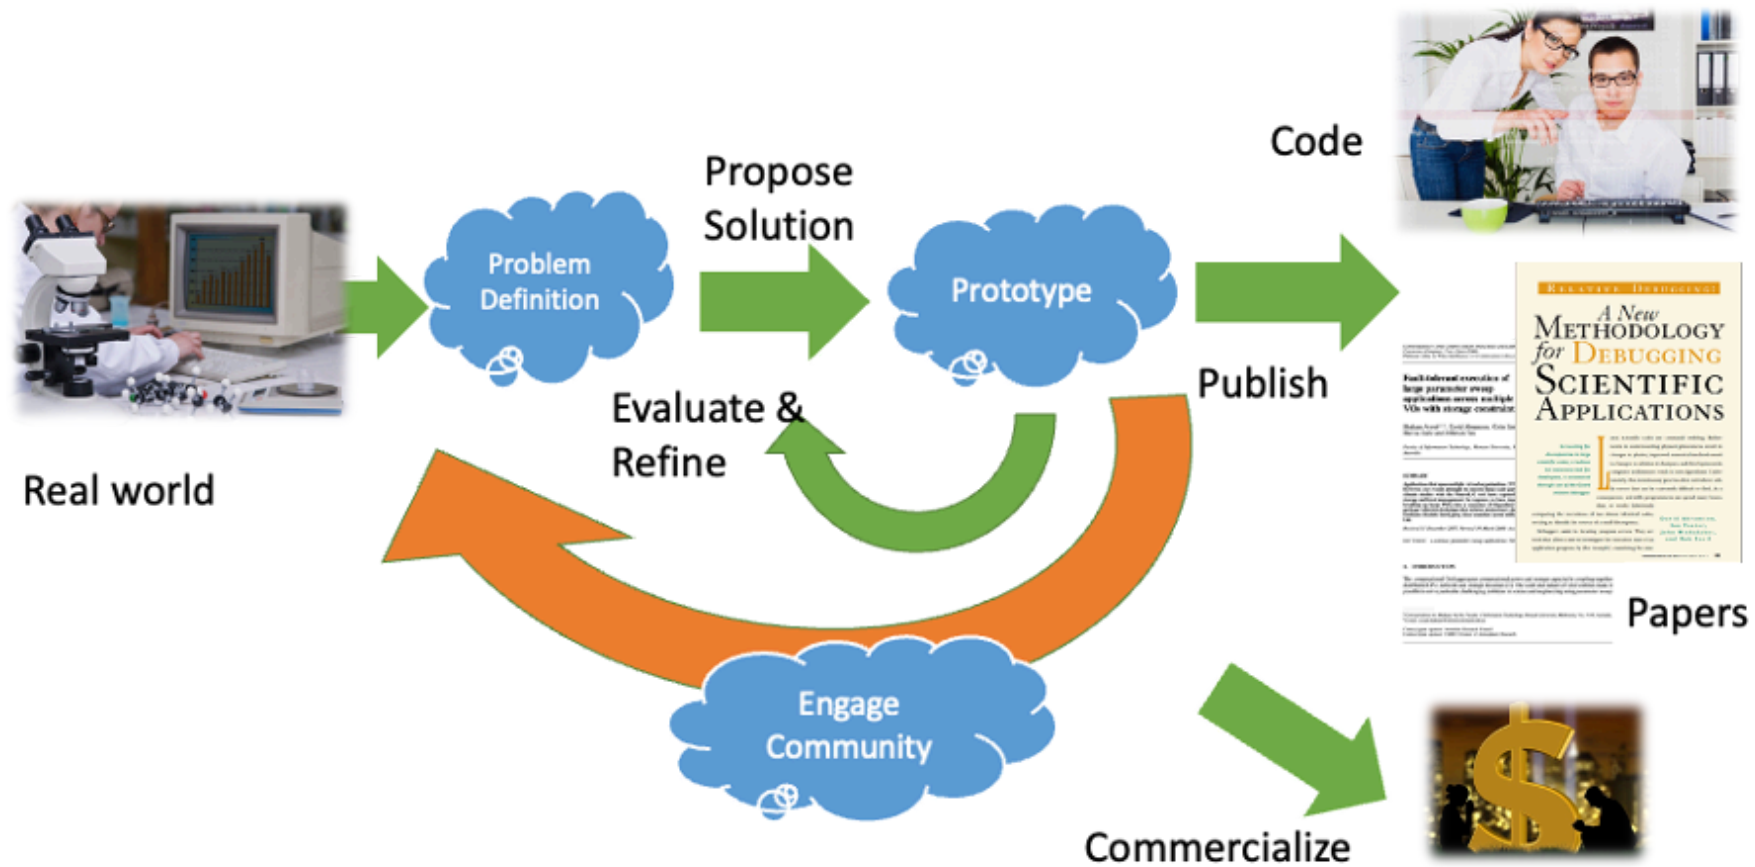
1. Provenance is useful for computational scientists
 - Record/replay, digital notebook, prov diff, create Nix package
2. Current state-of-the-art not practical use
 - must not modify kernel; must be rootless; must not crash; must be fast
3. But there is hope
 - lib. interposition

PROBE: Provenance for Replay OBservaTion Engine

<https://github.com/charmoniumQ/PROBE>



Difficulty of translational CS



Invitation for collaboration

- Using record/replay (*)
- Using provenance to generate Guix/Nix/Spack/Makeflow (*)
- Consuming provenance in other ways (WRROC *, *)
- Using the benchmark set
- sam@samgrayson.me
- <https://github.com/charmoniumQ/PROBE>

